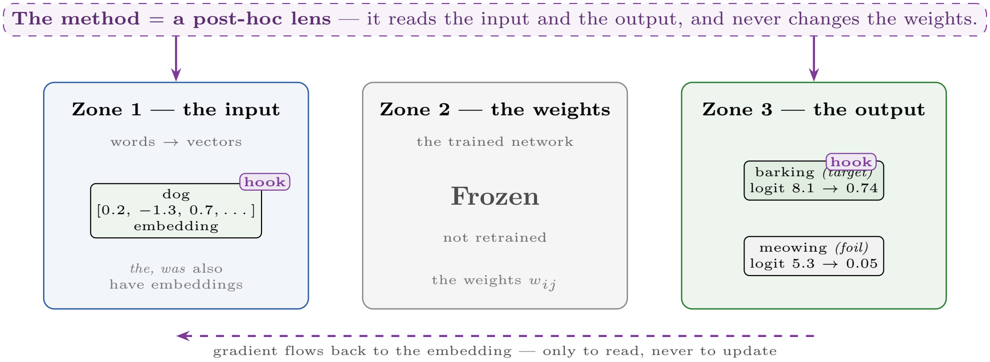

The method = a post-hoc lens — it reads the input and the output, and never changes the weights.
Zone 1 — the input
words → vectors
hook
dog
[0.2, −1.3, 0.7, . . .]
embedding
the, was also have embeddings
Zone 2 — the weights
the trained network
Frozen
not retrained
the weights w_{ij}
Zone 3 — the output
hook
barking *(target)*
logit 8.1 → 0.74
meowing *(foil)*
logit 5.3 → 0.05
gradient flows back to the embedding — only to read, never to update